

Introduction of 3D Vision

Hao Fang

方昊

深圳大学计算机与软件学院

可视计算研究中心 (VCC) : <https://vcc.tech/index.html>

地址：深圳大学沧海校区致真楼812

B.Sc: Beihang University

M.Sc: Beihang University

Engineer Program: Ecole Centrale Paris

Ph.D: Inria Sophia-Antipolis

Research of interest : 3D Vision

Email: paulfanghao@gmail.com

Personal web: <https://vcc.tech/~haofang>




Lecture 1

Introduction

Computer Vision

“One picture is worth ten thousand words”

Human beings



Machines



Computer Graphics

“Using machines to model / animate / render **virtual** objects”

Computer-Aided-Design




Computer Graphics = Modeling / Animating / Rendering “3D data” with “Algorithm”

3D Vision

“Using machines to understand or generate **real-world** models”

Reverse Engineering



3D Vision = Analyzing “3D data” with “Algorithm”

Prerequisites

Mathematics:

- Linear algebra: SVD, eigen values, eigen vectors, PCA...
- Geometry: point, 2D line, 3D plane, 2D polygon, intersections...
- Probabilistic & Statistic: Bayes, Maximum Likelihood, PDF...
- Optimization: convex, non-linear, discrete, IP, MRF, MCMC...

Coding:

- Data structure and algorithm
- C++: class, shared pointer, Polymorphisn...
- Python

Prerequisites

Computer Vision

- Basic Image Processing: SIFT, edge detection, image smoothing...
- High-level tasks: object detection, semantic segmentation...
- OpenCV

Deep Learning:

- Basic knowledge: loss function, networks, data preprocessing...
- Some tricks for training: batch training, regularization...
- PyTorch, Tensorflow

What you can learn from this course

The whole pipeline of **3D Vision**

- 3D data representation
- 3D data reconstruction
- 3D data processing
- 3D data understanding
- Real-world applications and methods
- Useful tools:
 - CGAL: mesh processing algorithms (C++)
 - PCL: point cloud processing algorithms (C++)
 - Open3D, PyTorch3D (Python)

Schedule

Lecture: 70% (me)

Paper reading and sharing: 30% (you)

Assignment: to be continued

Methodology


Whatever you learn, ask yourself three questions:

- Why? (the objective, potential applications...)
- What? (input, output, requirements...)
- How? (method, pipeline, technical details...)


Part 1

3D Data and Applications

2D vs 3D



projection




depth is not determined



It's non-trivial to infer 3D information from 2D images, so 3D data is more appropriate for real-world applications!

Computer Graphics & 3D Vision

Computer Graphics



Creating “virtual 3D model”

Rendering

3D Visison




Generating “real-world 3D model”

Understanding

3D Data

□ Voxel



$H * W * D * 1$

Pros:


- Regular: 3D grids
- Directly applying 3D CNN

Cons:

- Trade-off between resolution and computational cost

3D Data

❑ Point Cloud



$N \times C$ ($C=3$ or 4)

Pros:


- Easy to collect: 3D scanners
- High accuracy w.r.t the surface of objects

Cons:

- Unordered: no direct deep feature extrators
- No topology between local points
- Artifacts: outliers, noise, holes...
- Cannot be used for rendering
- Difficult for storing and manipulating

3D Data

❑ Polygonal Mesh



Vertices, Edges, Facets

Pros:

- Explicit topology information
- Used for Rendering
- Easy for manipulating and storing

Cons:

- Cannot be scanned directly: Sketchup, **Reverse Engineering**


Sensors

RGB Cameras

- **monocular**
- stereo
- video
- panoramic



$H * W * 3$




$$Z \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} = \begin{pmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} X \\ Y \\ Z \end{pmatrix}$$

Z (depth) is not determined

Sensors

RGB Cameras

- monocular
- **stereo**
- video
- panoramic




$$z = \frac{fb}{d}, \quad d = u_L - u_R$$

Z (depth) is determined by physical relation


Sensors

RGB Cameras

- monocular
- stereo
- **video**
- panoramic



$$T * H * W * 3$$



Z (depth) is determined by algorithms:
epipolar geometry and triangulation

Sensors

RGB Cameras

- monocular
- stereo
- video
- **panoramic**





Z (depth) is not determined

Sensors

RGB-D (depth) cameras

- Structured Light: Kinect-1
- Time-of-flight: Kinect-2





Z (depth) is determined

相机根据返回的结构光图案,计算物体离自身的距离

Sensors

RGB-D (depth) cameras

- Structured Light: Kinect-1
- Time-of-flight: Kinect-2




Z (depth) is determined

相机根据发送到返回之间的光束飞行时间,确定物体离自身的距离

Sensors

Laser scanners: Lidar

- 机械激光雷达
- 固态激光雷达




- expensive
- large size
- accurate

Sensors

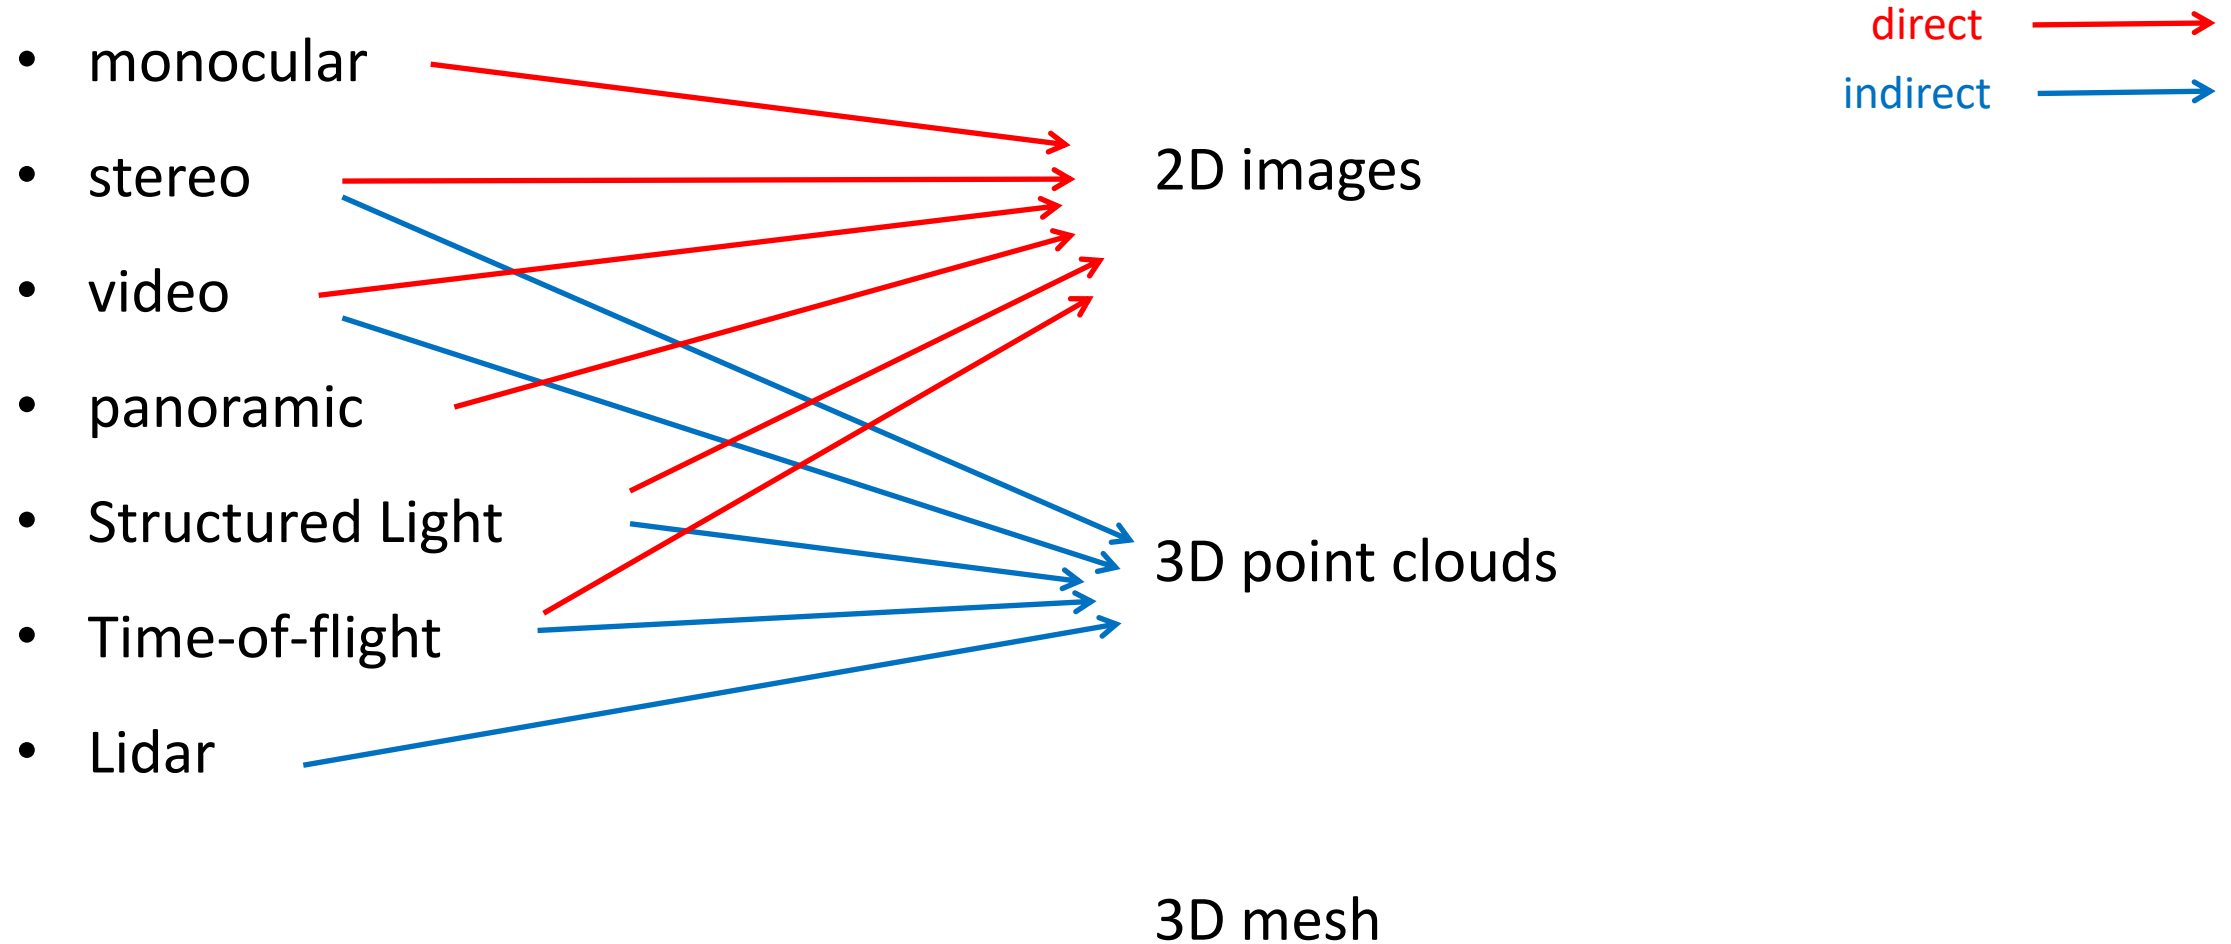
Laser scanners: Lidar

- 机械激光雷达
- 固态激光雷达



- cheap
- small size
- less accurate

Summary



Applications

Smart City / Digital Twin

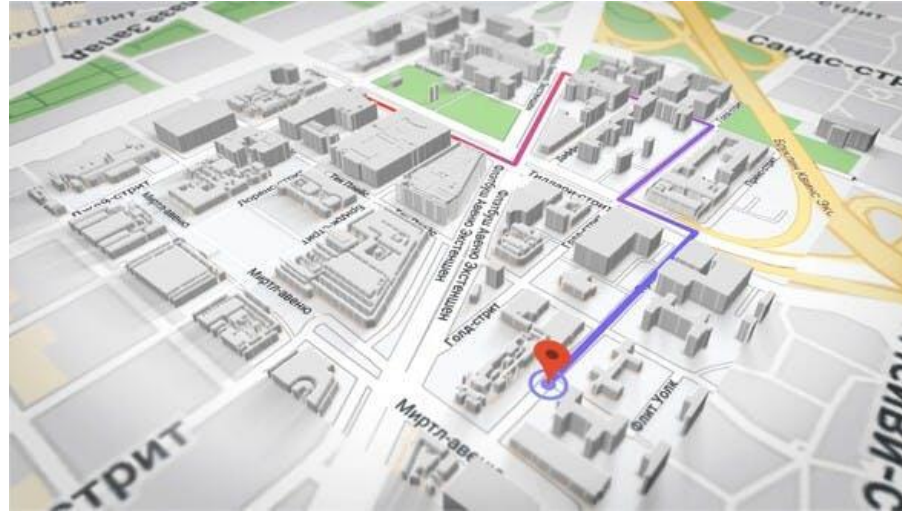
- urban planning
- mapping and navigation
- simulation and modeling
- security



Applications

Smart City / Digital Twin


- urban planning
- mapping and navigation
- simulation and modeling
- security



Applications

Smart City / Digital Twin

- urban planning
- mapping and navigation
- simulation and modeling
- security



Applications

Smart City / Digital Twin

- urban planning
- mapping and navigation
- simulation and modeling
- security



Applications

Smart Home


- navigation
- AI furnishing
- VR / AR



Applications

Smart Home

- navigation
- **AI furnishing**
- VR / AR



Applications

Smart Home

- navigation
- AI furnishing
- VR / AR



Applications

Industry

- autonomous driving / robotics
- packing & transport
- design



Applications

Industry

- autonomous driving / robotics
- packing & transport
- design



Applications

Industry

- autonomous driving / robotics
- packing & transport
- design



Applications

Intertainment

- animation
- games
- movie



Applications

Intertainment

- animation
- games
- movie



Applications

Intertainment

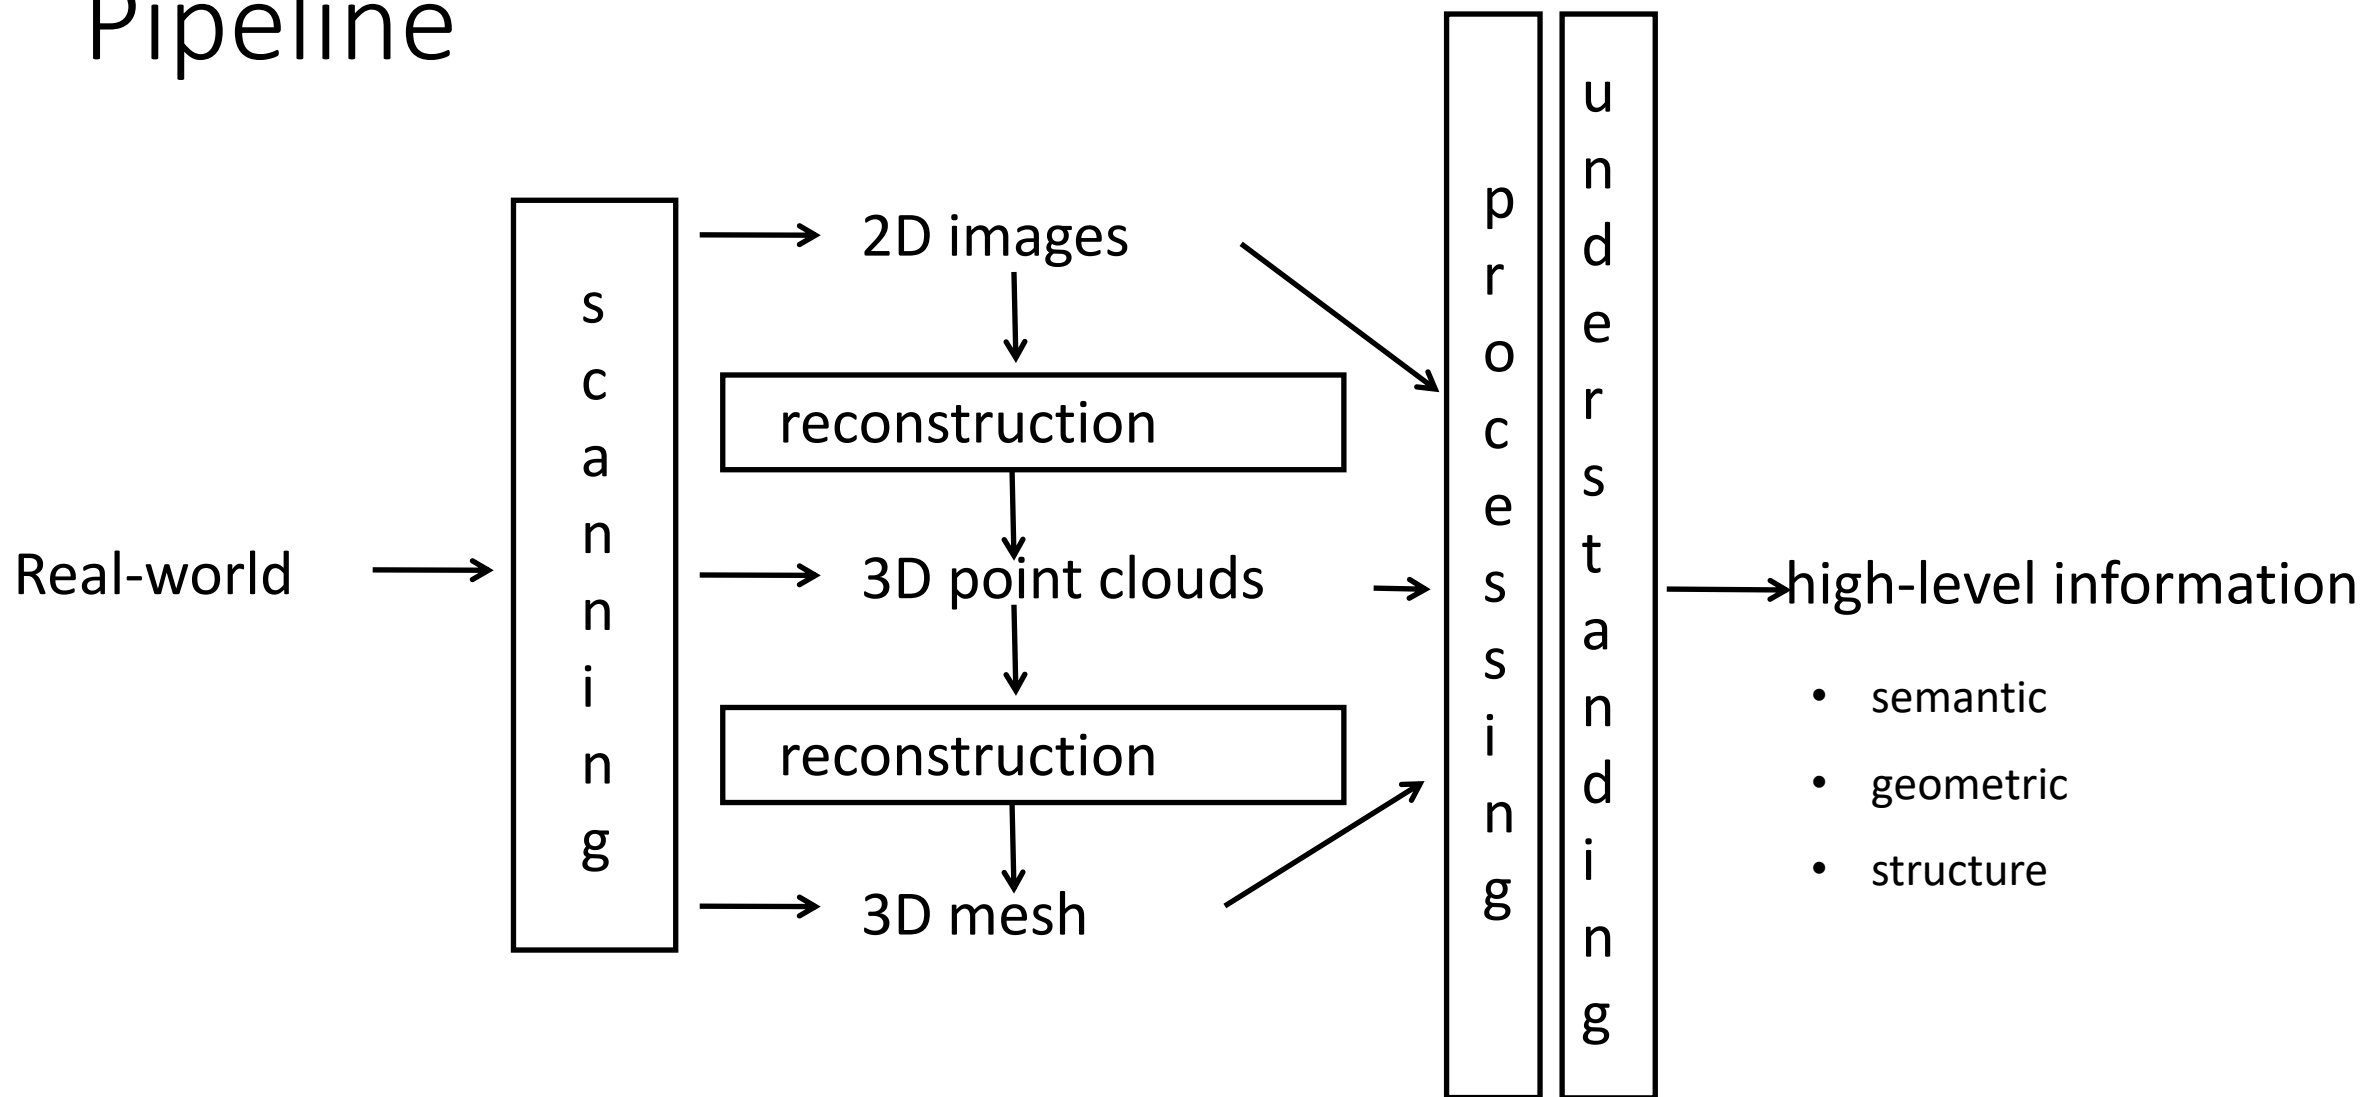
- animation
- games
- movie



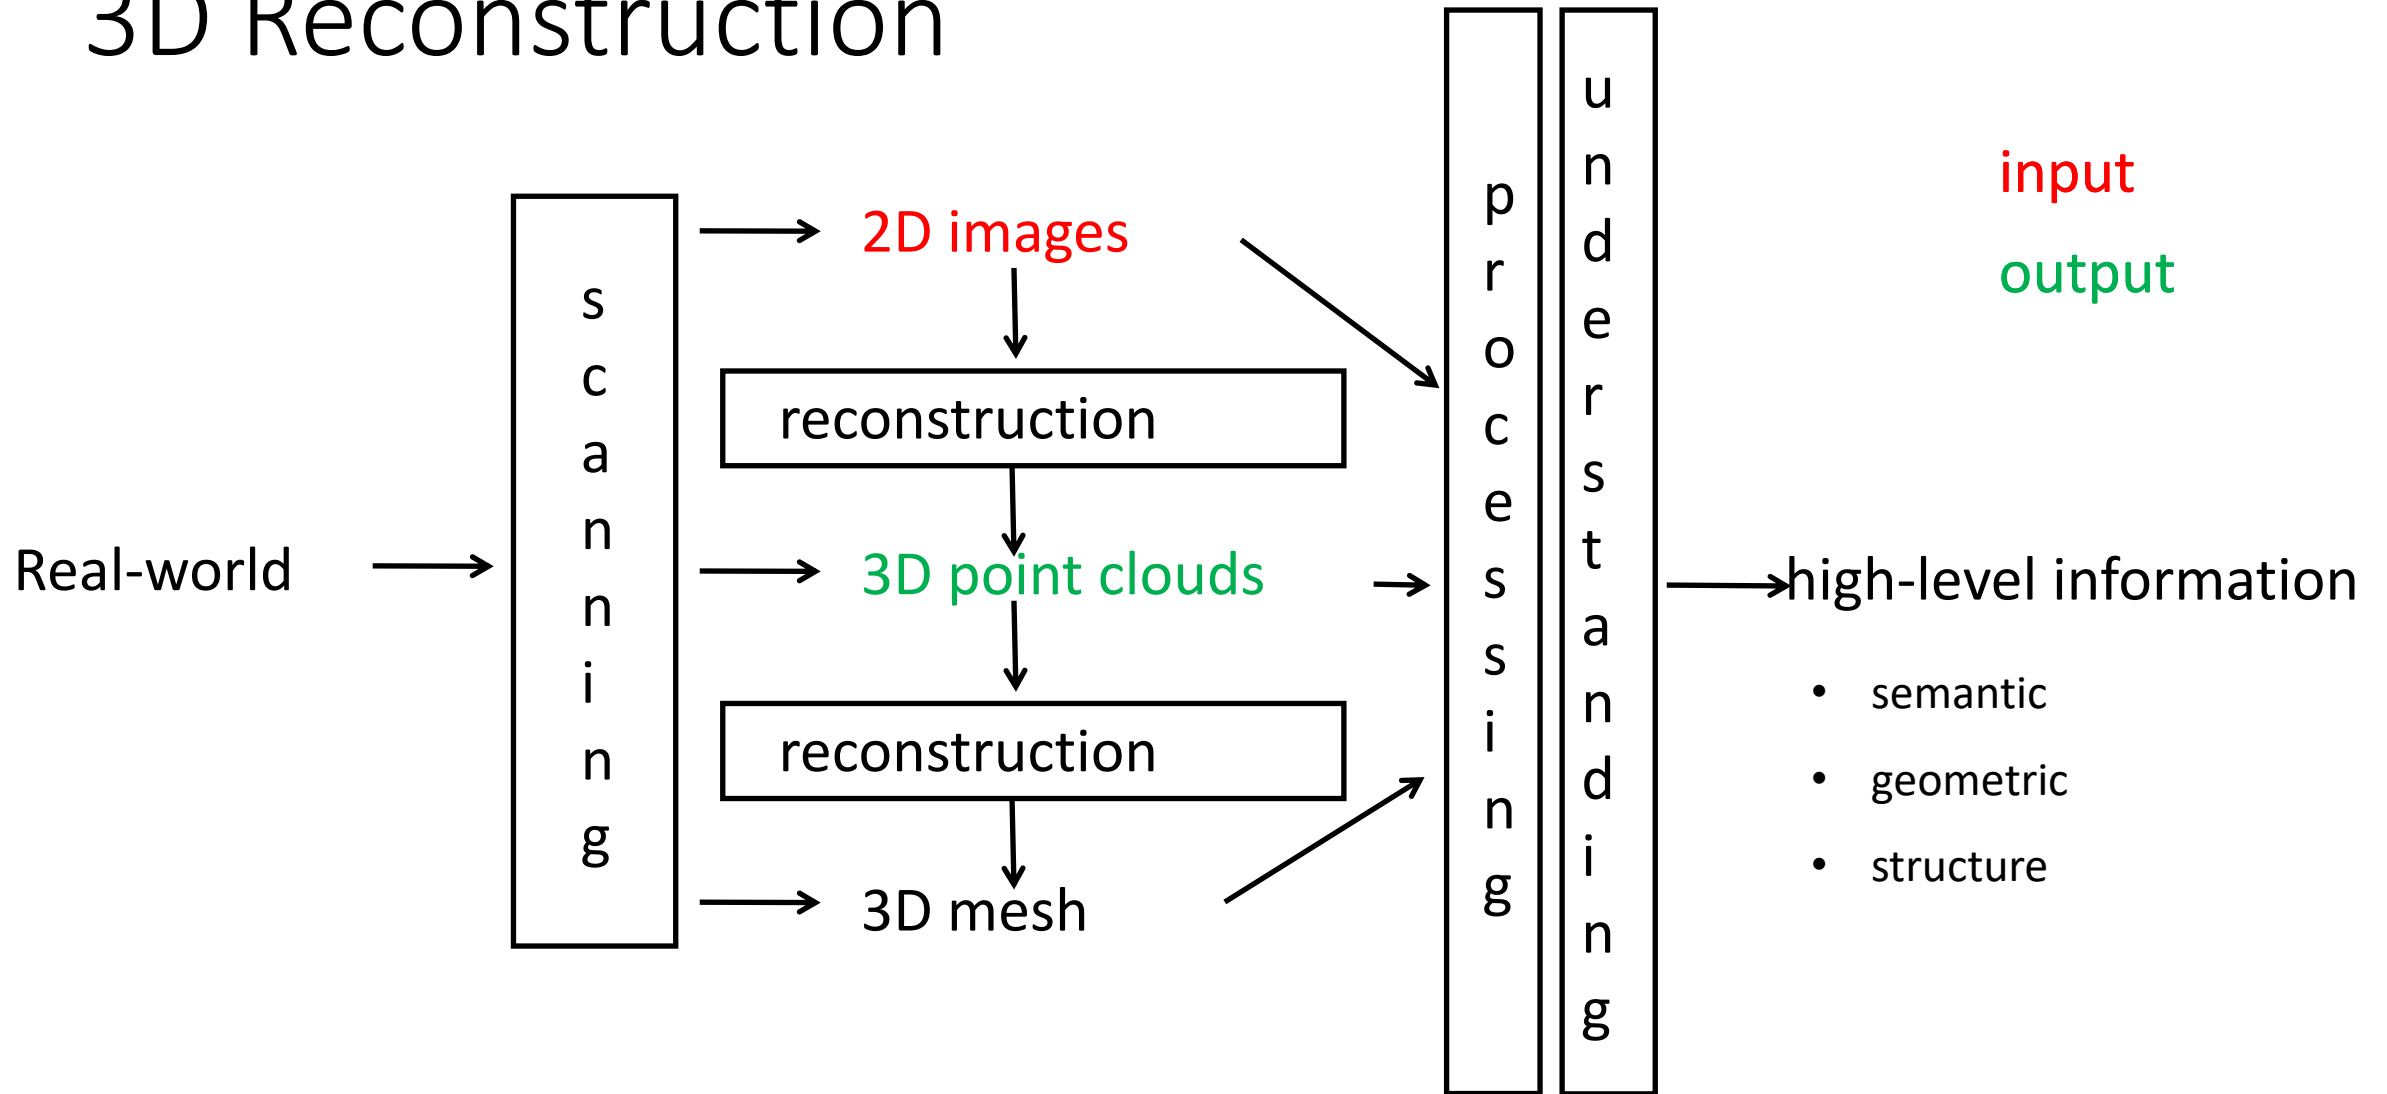
Part 2

3D Vision Pipeline

Pipeline



3D Reconstruction



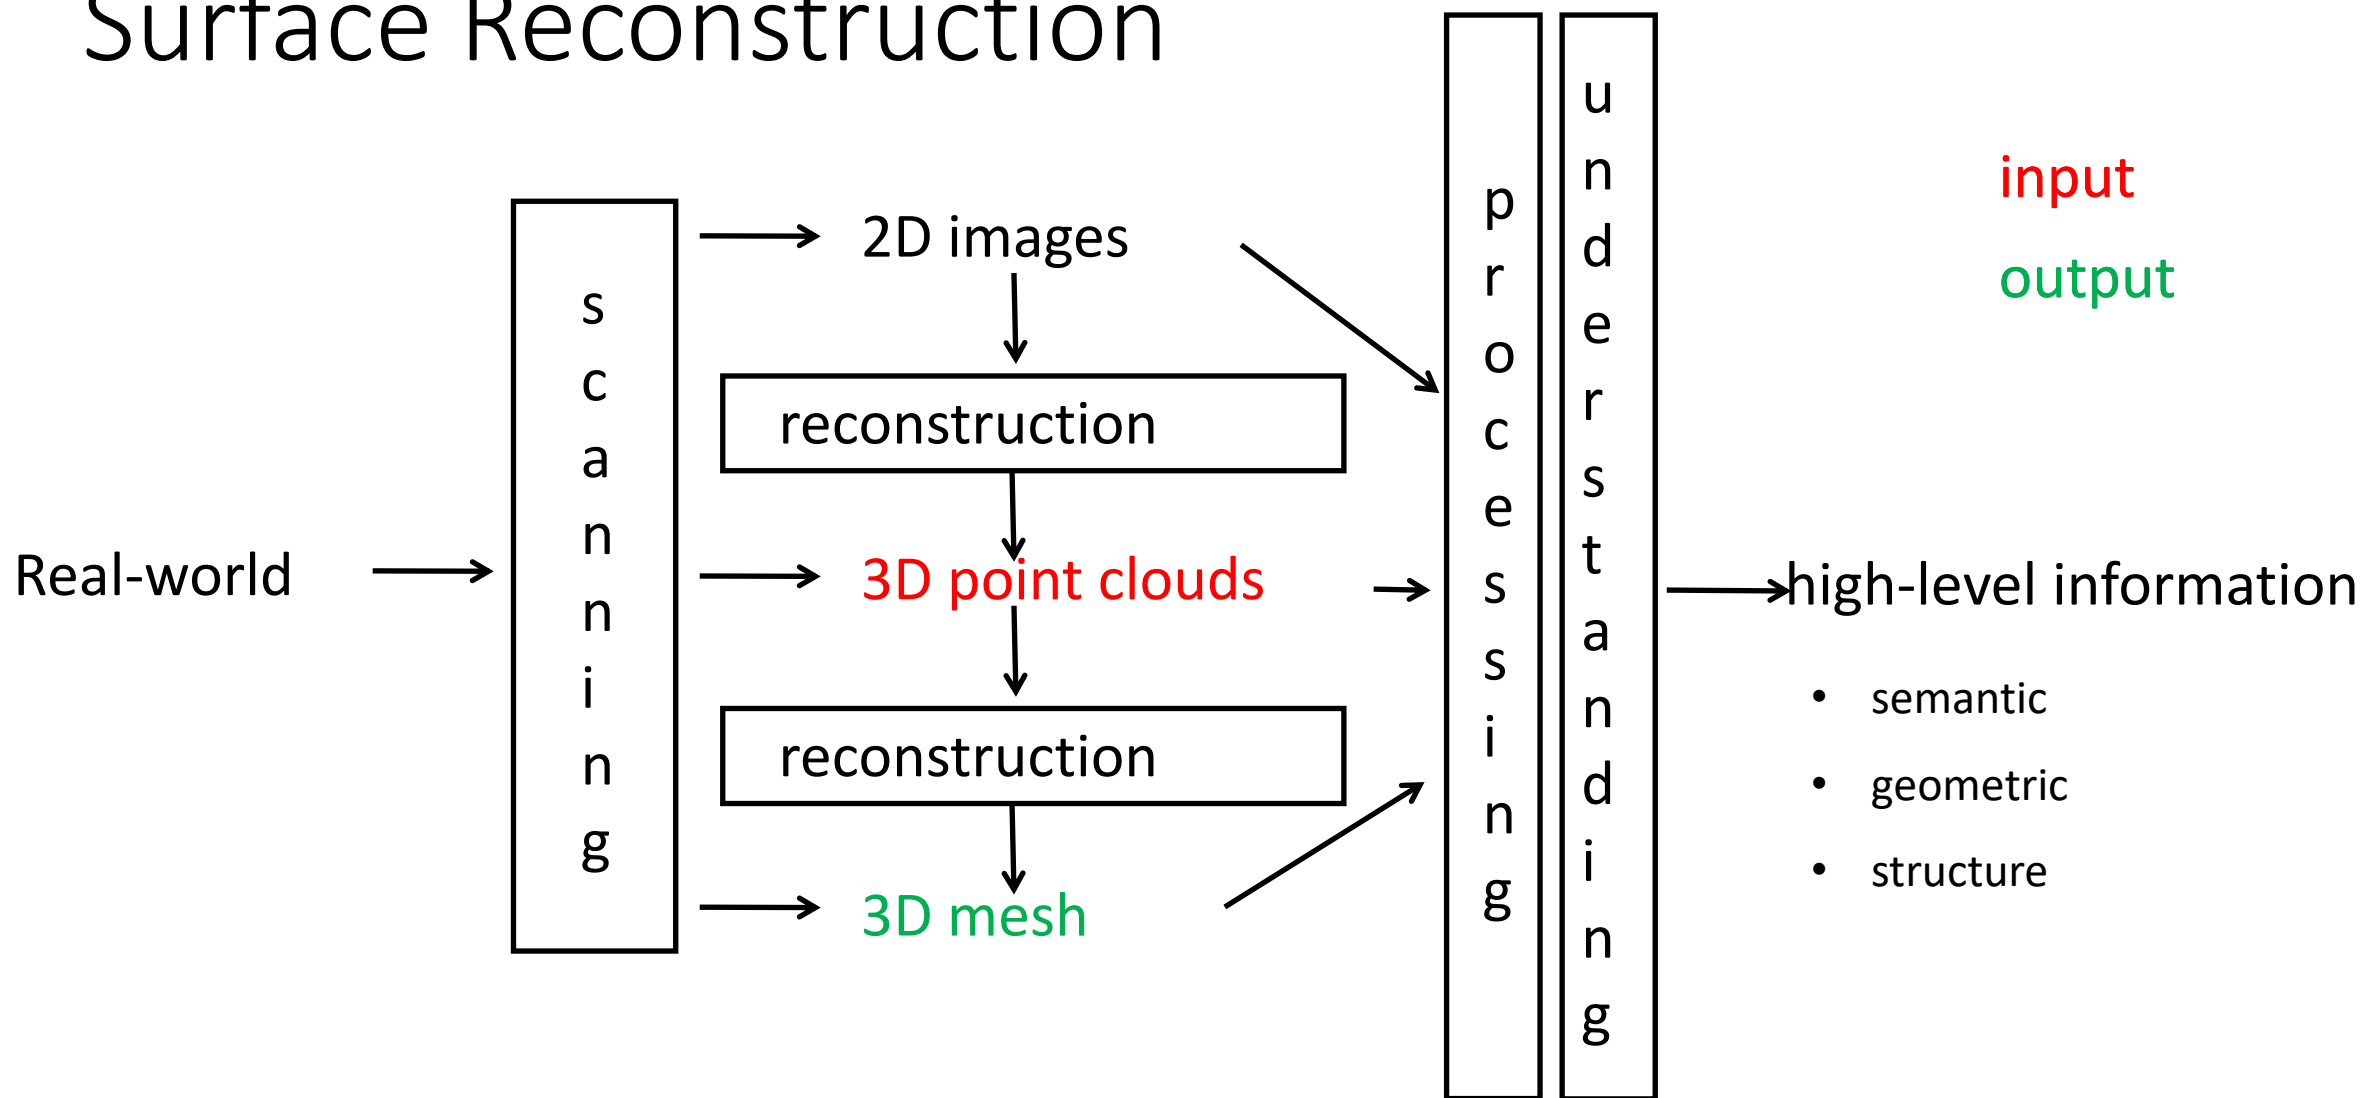
3D Reconstruction




input: 2D images

output: 3D point clouds

Surface Reconstruction



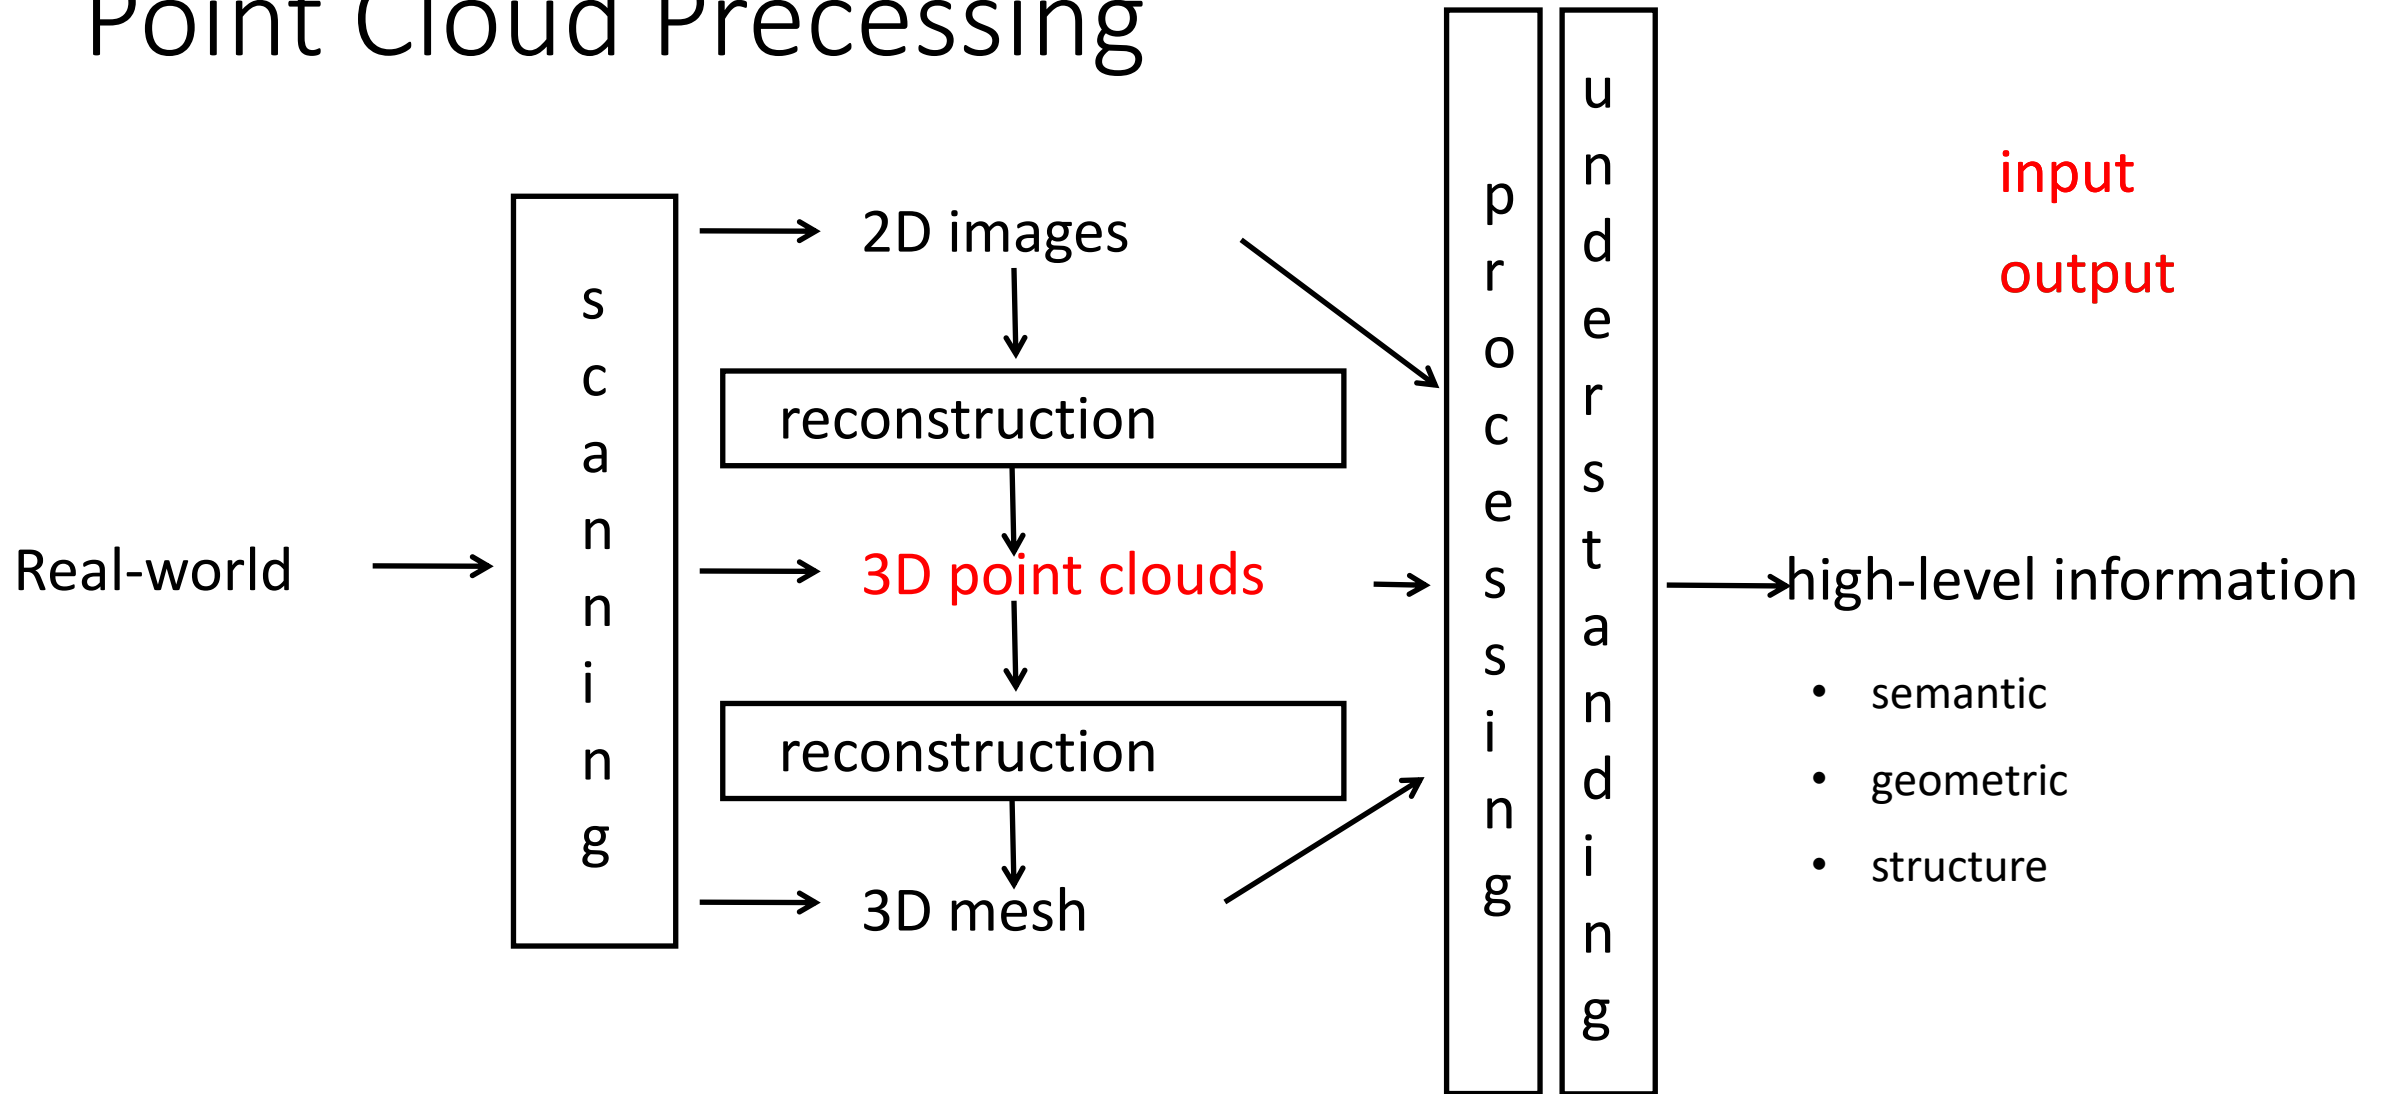
Surface Reconstruction



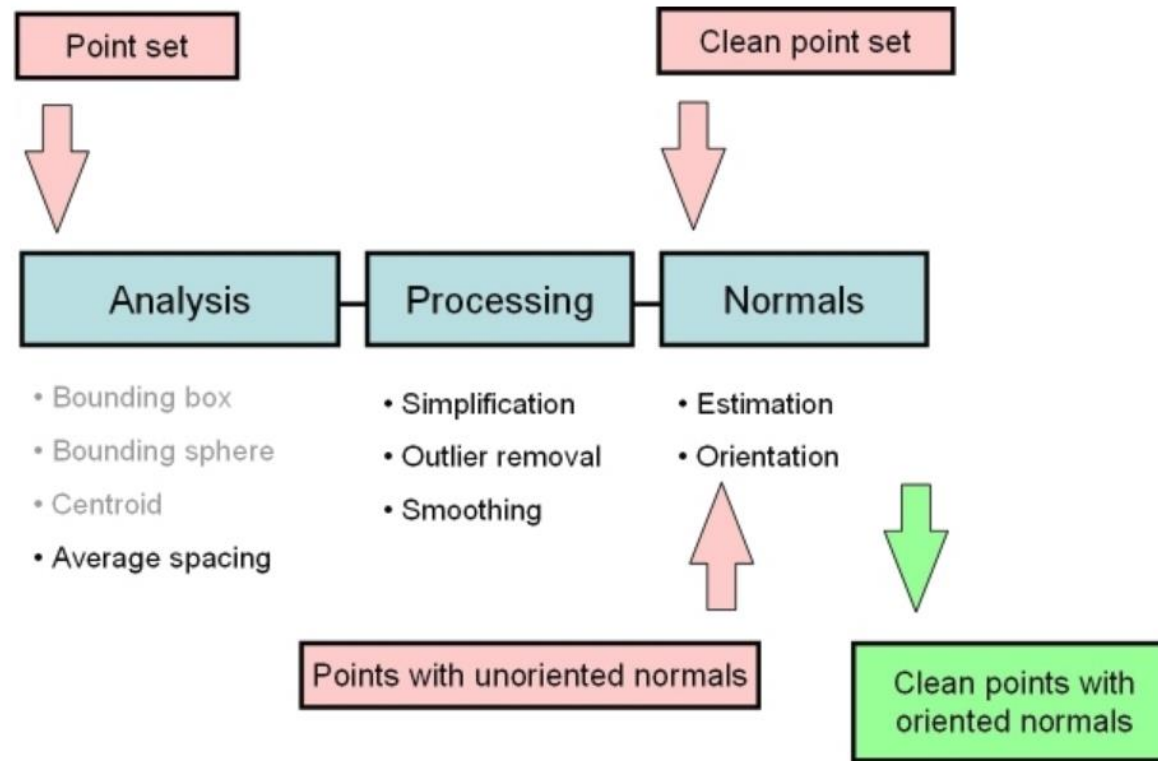
input: 3D point clouds

output: 3D mesh

Point Cloud Preprocessing



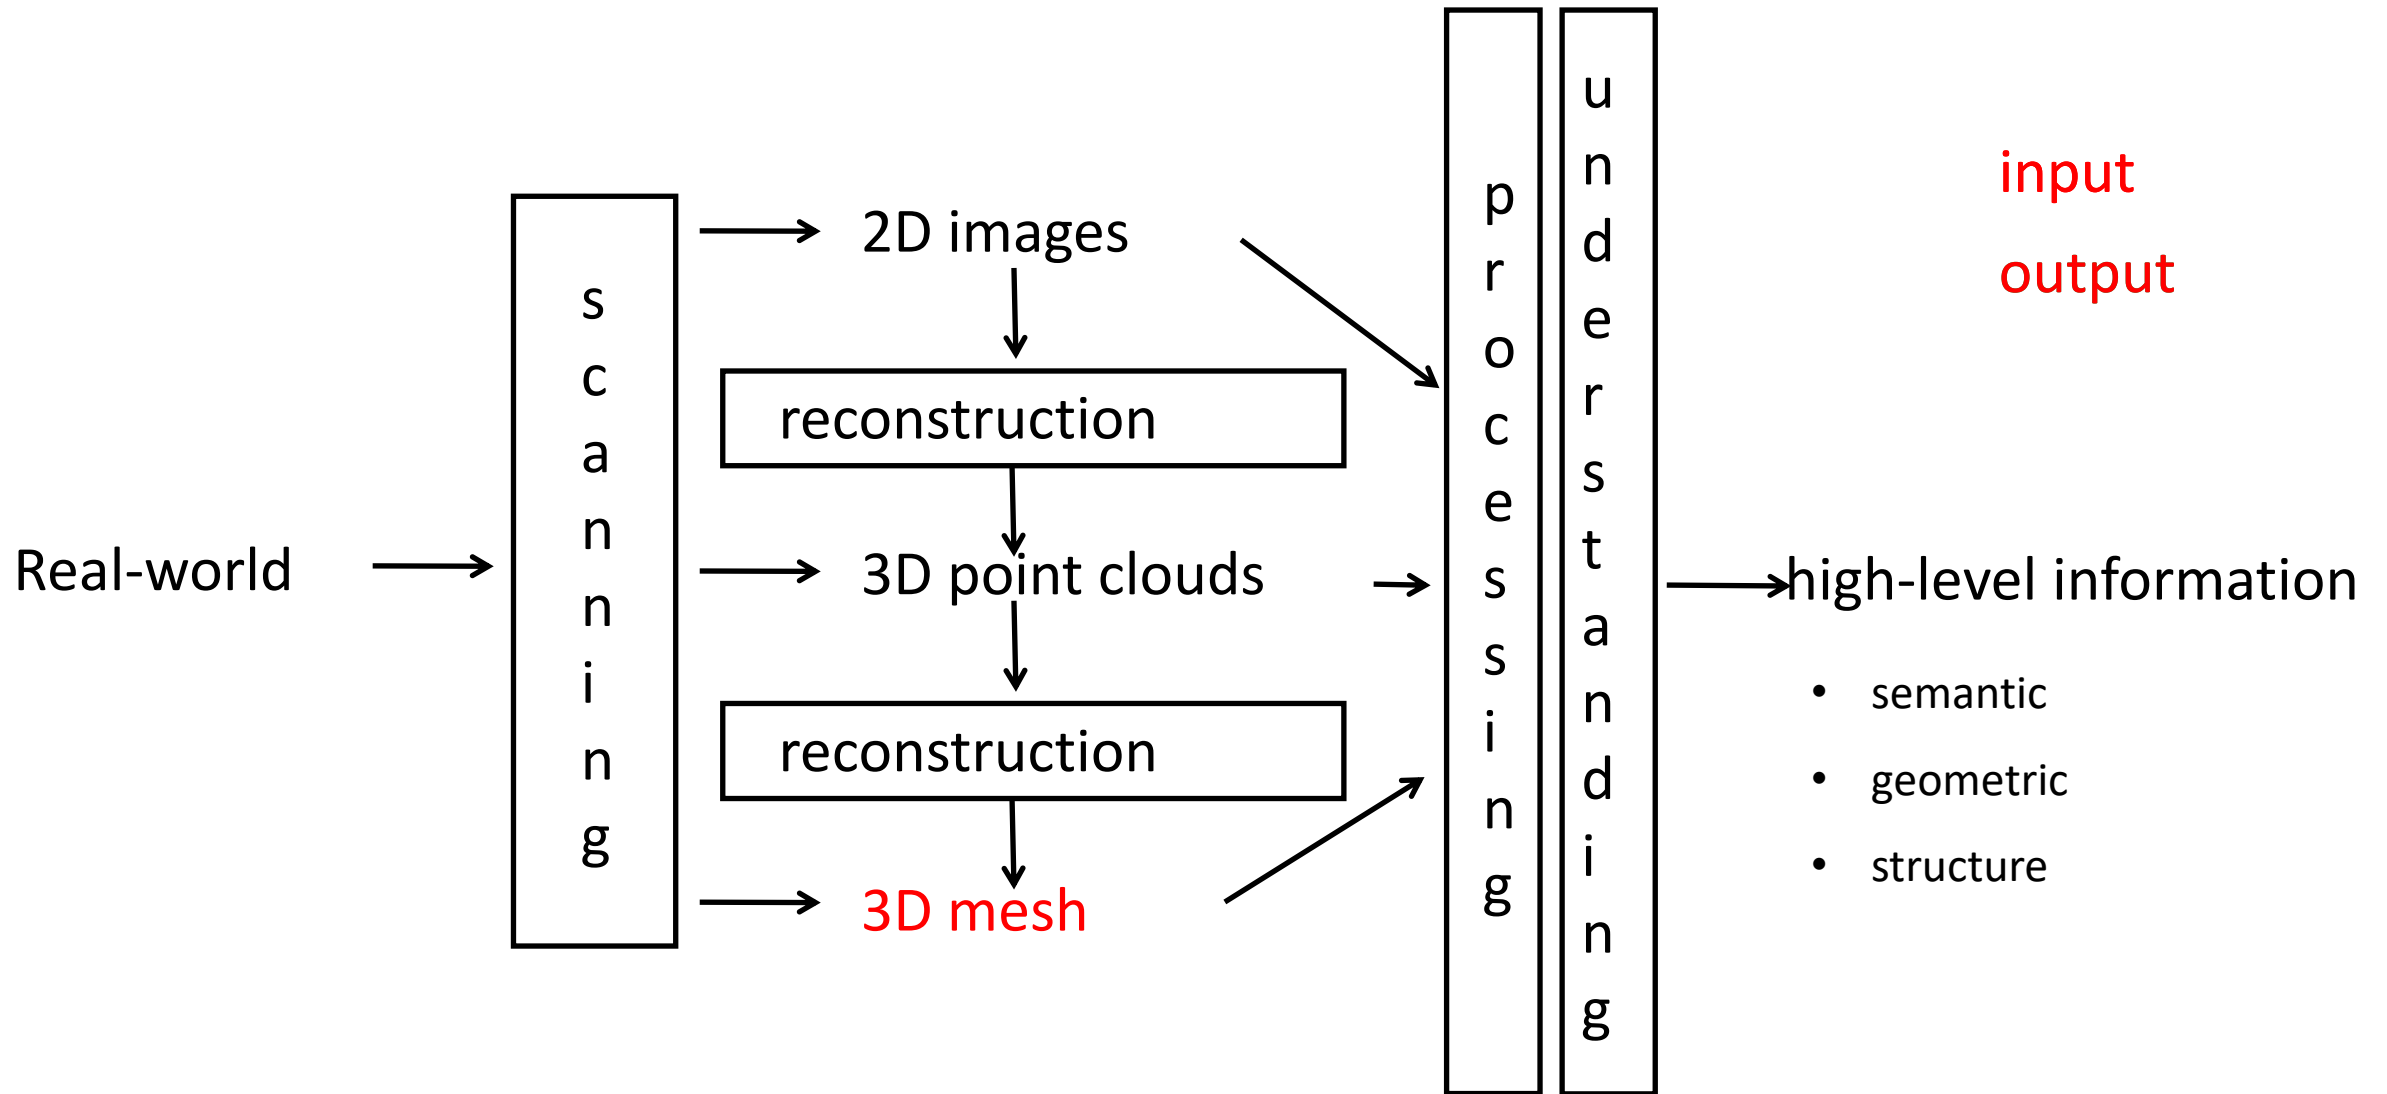
Point Cloud Preprocessing



input: raw point clouds


output: clean point clouds with normals

Polygonal Mesh Preprocessing



Polygonal Mesh Preprocessing

Smoothing



input: raw triangular mesh

output: smooth mesh

Polygonal Mesh Preprocessing

Hole filling





Figure 64.10 Results of the main steps of the algorithm. From left to right: (a) the hole, (b) the hole after its triangulation, (c) after triangulation and refinement, (d) after triangulation, refinement and fairing.

input: raw triangular mesh with holes

output: watertight mesh

Polygonal Mesh Preprocessing

Deformation




input: raw triangular mesh

output: defromed mesh

Polygonal Mesh Preprocessing

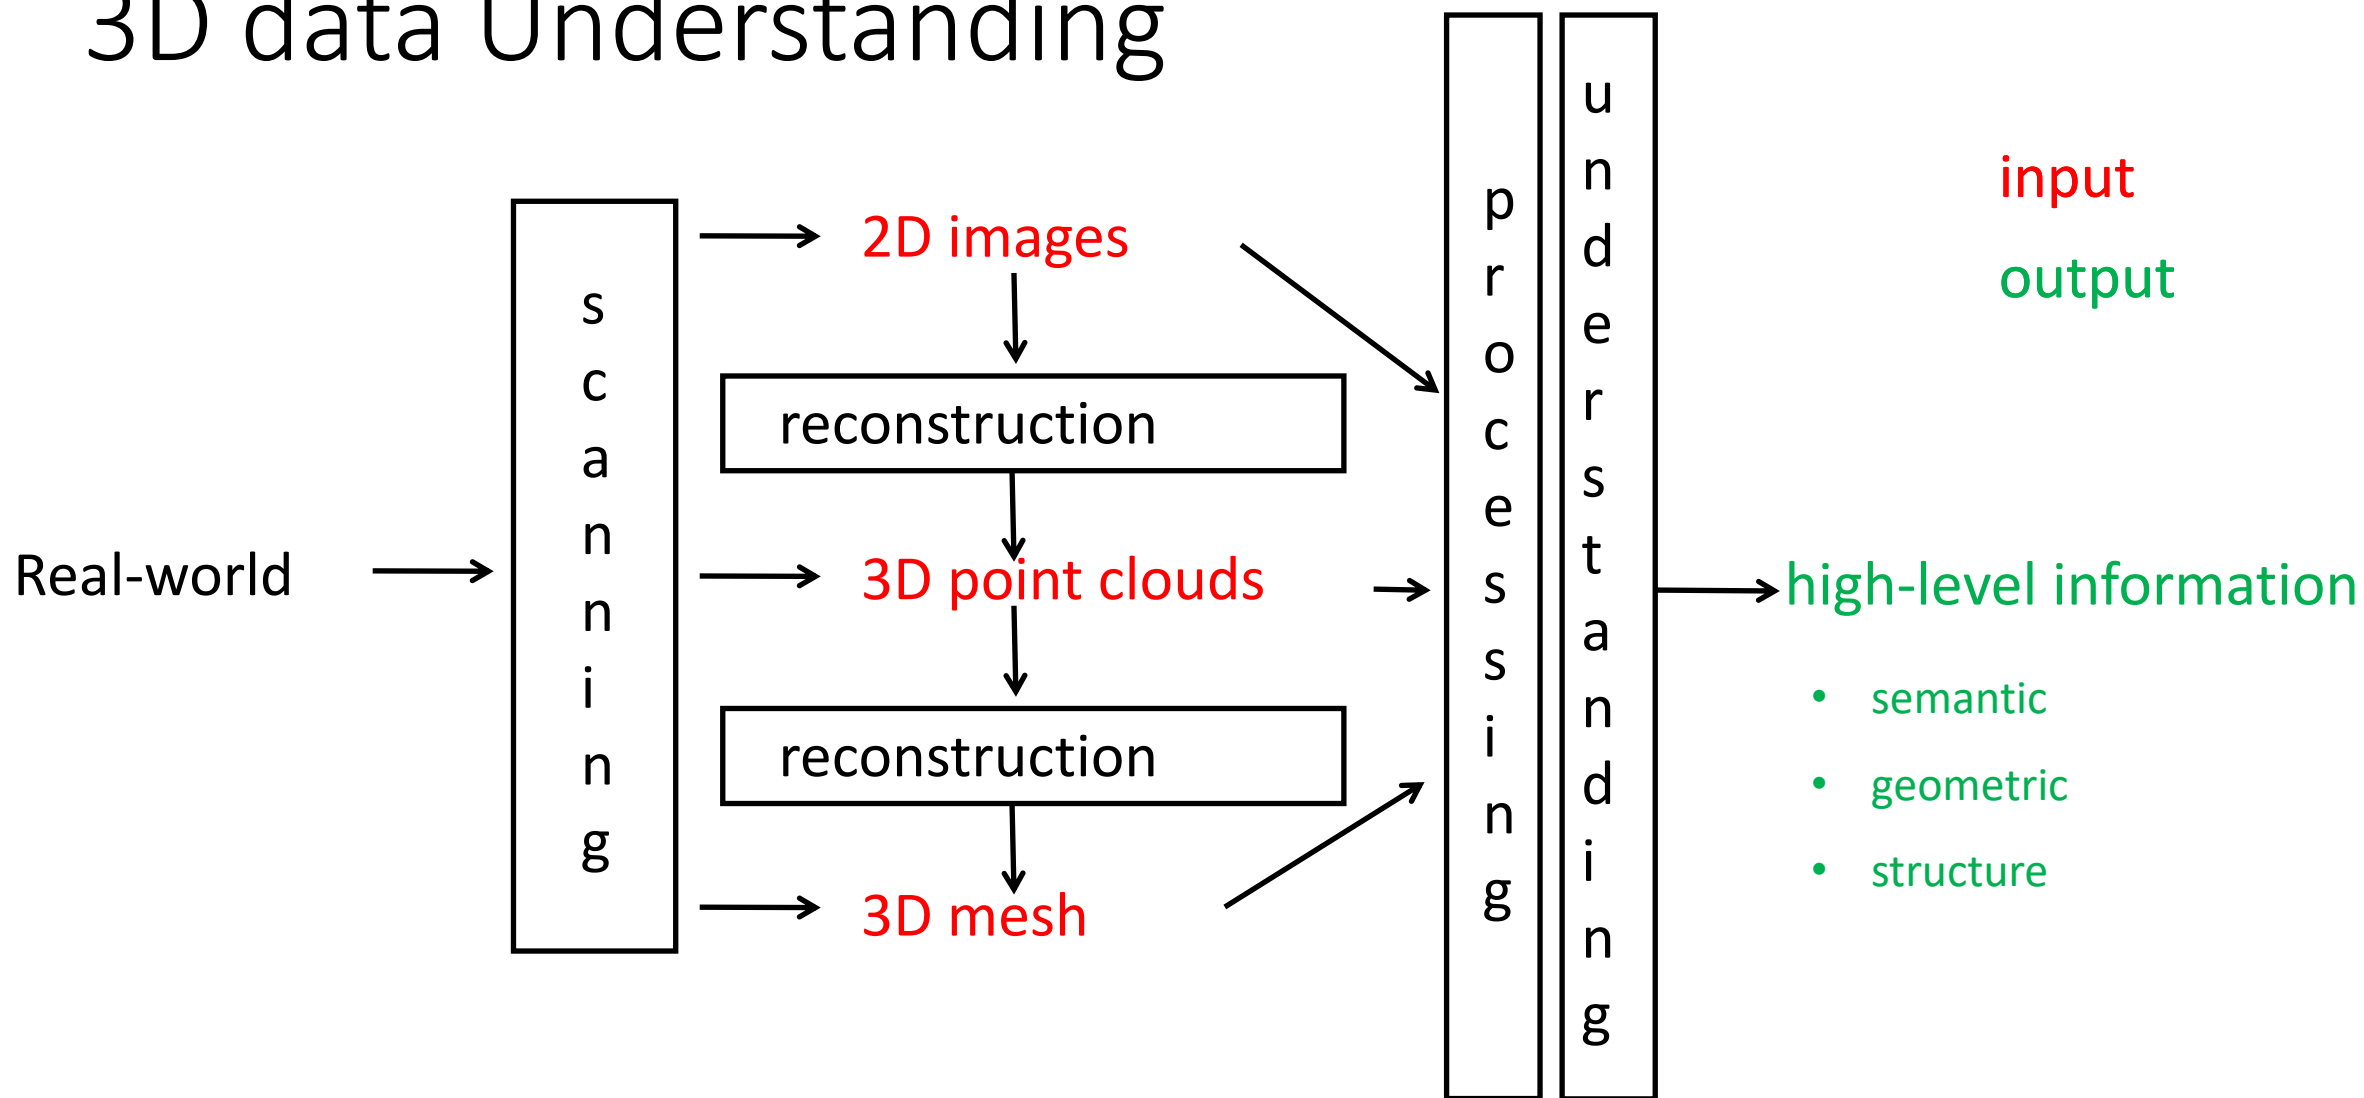
Simplification



input: raw triangular mesh



output: simplified mesh

3D data Understanding




3D data Understanding

3D semantic segmentation




raw point cloud or
triangular mesh



semantic label

3D data Understanding

3D object detection




input: point clouds


output: 3D bounding box

3D data Understanding

3D registration




input: two point clouds



output: transformation matrix

3D data Understanding

3D shape detection



input: point clouds or mesh

output: geometric shapes

3D data Understanding

scene / model completion




input: partial 3D data


output: complete 3D data

3D data Understanding

Floorplan generation



input: RGB images



output: 2D planar graph

3D data Understanding

Layout prediction

estimated layout (orange lines) compared with ground truth (green lines)



Figure 6. Qualitative results for non-cuboid layout prediction

input: panoramic image

output: scene layout

3D data Understanding

Depth prediction




input: video images

output: depth map

3D data Understanding

View Synthesis




input: RGB images

output: radiance at (x,y,z) from view direction (θ, φ)

Part 3

Important Tools

CGAL



Computational Geometry Algorithms Library

- 基本的几何数据结构：点，线，面，向量，多边形，多面体，点云，mesh等
- 基本的几何元素之间关系的算法：查询是否相交，计算距离
- 基本的点云，mesh数据处理算法
- 基本的优化算法


CGAL

- 可移植性高
- generic programming : template, STL, boost
- 处理浮点数计算的鲁棒性高
- 准确性高
- 效率高


CGAL

- 丰富的几何相关数据结构和算法：

- 2D :




Triangulations





2D Mesh Generation




Polyline Simplification




Visibility





Voronoi Diagrams



Arrangement



Boolean Operations




PCA



Neighbor Search




Minkowski Sum




Straight Skeleton


- 3D :




Triangulations
Voronoi Diagrams



Polyhedral Surface



Mesh Generation



Parameterization




Deformation



Boolean Operations




Convex Hull



Surface Reconstruction

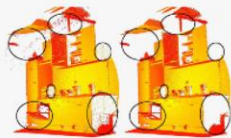
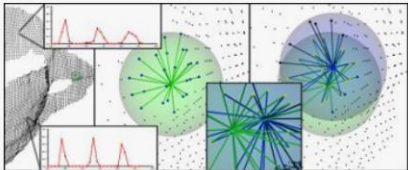
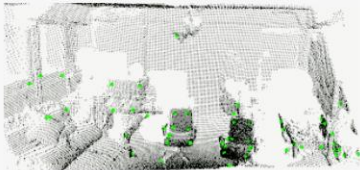
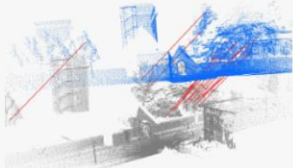
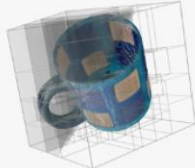


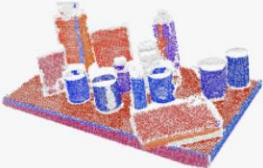
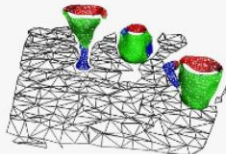
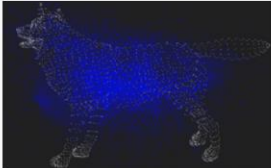




Simplification



Classification

PCL

filters	features	keypoints
		
registration	kdtree	octree
		
segmentation	sample_consensus	surface
		
recognition	io	visualization
		



pointcloudlibrary

PyTorch



- Dataset and dataloader are easy to use
- Rich deep learning operators (differential->end-to-end), i.e. CNN, FC, dropout, cross entropy loss, L2 loss...
- Famous network architectures, i.e. ResNet50, LSTM
- C++\CUDA based operators extensions

Paper & Code & Data

- Structured3D, S3DIS, ScanNet, KITTI...
- arXiv: <https://arxiv.org/>
- Conference: CVPR, ICCV, ECCV, SIGGRAPH, SIGGRAPH Asia
- Journal: PAMI, IJCV, ISPRS Journal, TOG, TVCG
- Webpage of famous researchers or laboratories
- Github: released code

3D Vision Lab and Researchers

- The Stanford Geometric Computation Group: <https://geometry.stanford.edu/>
- Facebook AI Research: <https://ai.facebook.com/research/>
- Hao Su (UCSD): <https://cseweb.ucsd.edu/~haosu/>
- Computer Vision and Geometry lab (CVG) , ETH, <http://www.cvg.ethz.ch/index.php>
- Angela Dai: <https://www.3dunderstanding.org/>
- Andreas Geiger: <http://www.cvlibs.net/>
- Vladlen Koltun: <http://vladlen.info/>
- Thomas A. Funkhouser: <https://www.cs.princeton.edu/~funk/>
- Matthias Nießner: <https://niessnerlab.org/publications.html>

3D Vision Lab and Researchers

- Visual Computing Research Center (VCC) , SZU: <http://vcc.szu.edu.cn/index.html>
- Kun Zhou: <http://kunzhou.net/>
- Ligang Liu: <http://staff.ustc.edu.cn/~lgliu/>
- Baoquan Chen: <http://cfcs.pku.edu.cn/baoquan/>
- Kai Xu: <https://kevinkaixu.net/>
- Xiaowei Zhou: <http://xzhou.me/>
- Jingyi Yu: <http://www.yu-jingyi.com/cv/>

Part 4

Homework

Assignment 1

- 0. Learn CMake
- 1. Install PCL
- 2. Install CGAL
- 3. Compile two projects on BB and run the test code
- 4. Upload the results to BB

Questions?